

ODEs in Stan

StanCon 2018 Helsinki

August 26, 2018

Daniel Lee

daniel@generable.com



**Before we
begin**

Overall Structure

1. Write ODEs in Stan; simulate from ODE
2. Estimate parameters of ODE for a single observation
3. Estimate parameters of ODE for multiple observations
4. Parallelization

What you should be comfortable with

- ▶ Writing Stan programs, including functions
- ▶ Running Stan from interface of choice.
For today:
 - ▶ CmdStan v2.18.0: running Stan
 - ▶ RStan: exporting data, reading fits
- ▶ Important: familiarity with Stan program blocks and Stan types

What we don't have time to cover. =(

- ▶ Full tutorial in Bayesian inference
- ▶ Full tutorial in Stan
- ▶ Full tutorial in ODEs
- ▶ Expert-level debugging through ODEs

Some Notes

- ▶ Working example is simple: simple harmonic oscillator
 - ▶ Highlights mechanics of ordinary differential equations (ODEs) in Stan
 - ▶ There are more complications with complex examples.
Feel free to use your own example.
- ▶ To get the most out of this course:
 - ▶ Work together!
 - ▶ Ask questions.

Last thing... install CmdStan v2.18.0

- ▶ Download the latest CmdStan:
<https://github.com/stan-dev/cmdstan/releases/tag/v2.18.0>
- ▶ Unzip the directory. Then type:
 - ▶ `make build -j4`
(or however many cores)
- ▶ Also clone <https://github.com/generable/stan-ode-workshop>

**Write ODEs in
Stan**

What's an ODE?


Differential equation describing the derivatives of a function with respect to an independent variable.

- ▶ What does that mean?
 - ▶ There's an independent variable, \mathbf{t} .
Usually time, but could represent anything.
 - ▶ There's a (multivariate) independent variable, \mathbf{y} , that depends on \mathbf{t} .
We want $\mathbf{y}(\mathbf{t})$ for some set of \mathbf{t} , but we don't know how to compute it directly.
 - ▶ We have a function that describes the derivative of the dependent variable with respect to the independent variable:
$$\mathbf{f}'(\mathbf{t}) = d\mathbf{y}(\mathbf{t})/d\mathbf{t} = g(\mathbf{t}, \mathbf{x}, \theta)$$

where \mathbf{x} are data, θ are parameters

Example: simple harmonic oscillator

- ▶ Two states.
- ▶ Equilibrium position: $(0, 0)$
- ▶ Restoring force proportional to displacement (with friction)



Example: simple harmonic oscillator


► Two states. y_1, y_2 .

► $\frac{dy_1(t)}{dt} = y_2$

► $\frac{dy_2(t)}{dt} = -y_1 - \theta^* y_2$

► If we know the initial conditions, $y_1(t_0), y_2(t_0)$, at time t_0 , and θ ,

we can solve for times $t_0 < t_1 < t_2 < \dots$



Stan ODE function

- ▶ The Stan function must have this signature:

- ▶ `real[] ode_system(real t, real[] y, real[] theta,
real[] x_r, int[] x_i)`

- ▶ Simple harmonic oscillator:

```
real[] sho(real t, real[] y, real[] theta, real[] x_r, int[] x_i) {  
  real dy1_dt = y[2];  
  real dy2_dt = -y[1] - theta[1] * y[2];  
  return { dy1_dt, dy2_dt };  
}
```

$$\frac{dy_1(t)}{dt} = y_2$$

$$\frac{dy_2(t)}{dt} = -y_1 - \theta * y_2$$

Let's break it down

```
real[] sho(real t, real[] y, real[] theta, real[] x_r, int[] x_i) {  
    real dy1_dt = y[2];  
    real dy2_dt = -y[1] - theta[1] * y[2];  
    return { dy1_dt, dy2_dt };  
}
```

- ▶ Number of states: 2.
- ▶ Return: $dy1/dt$, $dy2/dt$, evaluated at time t with state y
- ▶ Arguments:
 - ▶ t , independent variable
 - ▶ θ , ODE arguments that depend on parameters
 - ▶ x_r , ODE arguments that depend only on data
 - ▶ x_i , integer ODE arguments that depend only on data

integrate_ode_bdf() and integrate_ode_rk45()

- ▶ These functions solve for the states at the times given.
They take the same arguments.
 - ▶ BDF: backward-differentiation formula for stiff systems
 - ▶ RK45: 4th and 5th order Runga-Kutta method for non-stiff systems
- ▶ `real[] integrate_ode_*(ode_function, real[] initial_states, real t0, real[] times, real[] theta, real[] x_r, int[] x_i)`
- ▶ The new arguments:
 - ▶ `ode_function`: `real[] f(real, real[], real[], int[])`
 - ▶ `real[] times`. Times to solve for.
 - ▶ `return`: states at the times specified

Generate data from this ODE

- ▶ From CmdStan:
 - ▶ `make ../stan-ode-workshop/sho_sim`
 - ▶ `make ../stan-ode-workshop/sho_sim.exe`
- ▶ From stan-ode-workshop
 - ▶ `./sho_sim sample`
- ▶ Read from R:
 - ▶ `library(rstan)`
 - ▶ `fit <- read_stan_csv("output.csv")`
 - ▶ `plot(extract(fit)$y[1,,])`
- ▶ why is every "row" of y the same?
- ▶ Exercise: adjust program so it accepts times as data
- ▶ Exercise: add noise in the generation process

Finish simulation

- ▶ Create a new variable, `y_hat`, with `normal(0, 0.5)` noise for both states.
- ▶ Generate dataset
 - ▶ `data <- list()`
 - ▶ `data$T <- 20`
 - ▶ `data$ts <- 1:20`
 - ▶ `data$to <- 0`
 - ▶ `data$yo <- c(1, 0)`
 - ▶ `data$y_hat <- y_hat[1,,]`
- ▶ `stan_rdump(ls(data), "sho.data.R", envir = list2env(data))`

Questions

- ▶ Can the ODE be discontinuous?
- ▶ Can a model have more than one differential equation?
- ▶ If I can solve the ODE analytically, should I?

**Estimating from a
single observation**

Write a Stan program to fit the data

- ▶ What should be estimated?
 - ▶ sigma
 - ▶ theta
 - ▶ y0?
- ▶ Set reasonable priors (or you'll get in trouble)
 - ▶ $\text{sigma} \sim \text{normal}(0.5, 0.1);$
 - ▶ $\text{theta} \sim \text{normal}(0, 0.5);$
 - ▶ $\text{y0} \sim \text{normal}(0, 1);$

Run

- ▶ `./sho_fit sample data file=sho.data.R output refresh=10`
- ▶ `fit <- read_stan_csv("output.csv")`
- ▶ `print(fit, c("sigma", "yo", "theta"))`

More specifics for `integrate_*`

- ▶ Times:
 - ▶ Must be sorted
 - ▶ $\text{times} > \text{to}$
 - ▶ times must be data only
- ▶ `x_r` must be data only
- ▶ `x_i` must be data only
- ▶ Three optional parameters
 - ▶ relative tolerance: default $1\text{e-}10$
 - ▶ absolute tolerance: default $1\text{e-}10$
 - ▶ `max_steps`: default $1\text{e}8$


Some more implementation details

- ▶ We've only written $\frac{dy}{dt}$
- ▶ We solve for $y(t)$, for some set of ts that we specify. That is, we get:
 - ▶ $y(ts[1]), y(ts[2]), \dots, y(ts[T])$
 - ▶ we've solved for the $y(t)$ numerically
- ▶ But wait... Stan needs gradients with respect to parameters!
 - ▶ We also need $\frac{dy(ts[1])}{d\theta}, \frac{dy(ts[2])}{d\theta}, \dots, \frac{dy(ts[T])}{d\theta}$
 - ▶ The ODE system we're actually solving for in Stan has this many states:
 - ▶ (# of states) x (# of thetas)
 - ▶ If we need to solve for initial conditions, we also need (# of thetas)

**Estimate from
multiple observations**

Let's update simulation

- ▶ We want 10 sets of observations
- ▶ Let's draw 10 thetas, then generate observations.
- ▶ What does the data look like?
 - ▶ `real y_hat[10, T, 2];`
- ▶ Generate data.
If we're behind, use `sho_multiple.data.R`



Exercise: extend sho_fit.stan to multiple

- ▶ Hints:

- ▶ You'll want a loop over each individual

- ▶ `real theta[10];`

- ▶ Run with `iter=0`

- ▶ `theta ~ normal(0.15, 0.1);`

Review

- ▶ ODE function is the same
 - ▶ simulation
 - ▶ fitting to one set of observations
 - ▶ fitting to multiple set of observations
- ▶ ODEs are part of the model
 - ▶ Statistical model built on top
 - ▶ Estimates of parameters of the model happen simultaneously

Parallelization

Overview

- ▶ We write a function to handle each observational unit separately
 - ▶ Key requirement: the function is able to handle an independent part
- ▶ We pass this function into the **map_rect(...)** function
 - ▶ This "maps" the arguments to the function
 - ▶ All data is rectangular, so we'll need to pad
- ▶ From the language, this is all that's necessary

- ▶ We need separate build instructions for CmdStan (or interface of choice)

The function we need to write

- ▶ (vector, vector, data real[], data int[]) : vector
 - ▶ 1st vector: shared parameters
 - ▶ 2nd vector: individual parameters (both of these would have been theta)
 - ▶ data real[]: real data
 - ▶ data int[]: int data
 - ▶ return: vector. We have the flexibility to return what we want.
- ▶ This function will handle each observation

map_rect

- ▶ `vector map_rect((vector, vector, real[], int[]):vector f,
vector phi,
vector[] thetas,
data real[,] x_rs,
data int[,] x_is)`
- ▶ Stan will loop over all the parameters.
- ▶ The output will just be concatenated into one vector.

Let's write the function for one observation

```
vector individual_ode(vector phi, vector theta,  
                     real[] x_r, int[] x_i) {  
    int T = x_i[1];  
    real y[T, 2] = integrate_ode_bdf(sho, x_r[1:2], 0.0, x_r[3:(T+2)],  
    to_array_1d(theta), x_r, x_i);  
    return append_row(to_vector(y[, 1]), to_vector(y[, 2]));  
}
```

- ▶ Still calling `integrate_ode_bdf()`
- ▶ Return is a single vector of `y[, 1]` and `y[, 2]` appended.

Make the single function work in a loop

- ▶ See: `sho_fit_multiple_parallel_start.stan`
- ▶ Important parts:

- ▶ Setting up data

```
transformed data {  
  int x_i[N] = rep_array(T, N);  
  real x_r[N, T + 2];  
  vector[0] phi;  
  
  for (n in 1:N) {  
    x_r[n, 1:2] = y0[n];  
    x_r[n, 3:(T + 2)] = ts;  
  }  
}
```

- ▶ Unpacking the return

```
transformed parameters {  
  real y[N, T, 2];  
  
  for (n in 1:N) {  
    vector[T * 2] result = individual_ode(phi, to_vector(theta[n]), x_r[n], x_i);  
    y[n, , 1] = to_array_1d(result[1:T]);  
    y[n, , 2] = to_array_1d(result[(T + 1) : (2 * T)]);  
  }  
}
```


Using map_rect

- ▶ See: `sho_fit_multiple_parallel.stan`
- ▶ Important parts
 - ▶ Single call to `map_rect()`:

```
transformed parameters {  
  real y[N, T, 2];  
  
  {  
    vector[N * T * 2] result = map_rect(individual_ode, phi, theta, x_r, x_i);  
    for (n in 1:N) {  
      int start = T * (n - 1);  
      y[n, , 1] = to_array_1d(result[(start + 1) : (start + T)]);  
      y[n, , 2] = to_array_1d(result[(start + T + 1) : (start + 2 * T)]);  
    }  
  }  
}
```

To enable threading

- ▶ We need to rebuild the Stan program. Threads is easier to set up under linux + Mac

- ▶ Open a file, make/local

```
CXXFLAGS += -DSTAN_THREADS
```

- ▶ Rebuild program

```
touch ../stan-ode-workshop/sho_fit_multiple_parallel.stan  
make ../stan-ode-workshop/sho_fit_multiple_parallel
```

- ▶ Set environment variable

```
export STAN_NUM_THREADS=-1
```

- ▶ Run! It'll run on multiple threads.

What we didn't cover

- ▶ Ragged data
- ▶ Censoring
- ▶ Speeding up the code
- ▶ Debugging through an ODE

Thank you